

```
// Read successor child from disk
```

```

if (z.n > t - 1)           //Check if the successor child has t -1 keys.
    k' = Find-Successor-Key (k, y) // Locate the successor k' of a given key k in a subtree rooted at
                                   y (smallest key less than k)
    B-Tree-Delete-Key(z,k') // Recursively delete k' from the successor child
    x.keyi=k'                // Replace k with k'
    Disk-Write(x)           //Write node x back to disk
    Return true             //Return

```

### Case 2c

```

if (z.n == t-1) && (y.n == t-1) // Check if the preceding as well as successor child has only t-1 keys.
    y.keyn+1 = k                // Insert key k into y
for tmp1 = 1 to z.n            //for each key in z
    y.keyn+1+ tmp1 = z.keytmp1 // Append key from z to the end of y
y.n = y.n + 1 + z.n           // Update the size of y to reflect the merged node (y and z)
for tmp2 = i + 1 to x.n - 1    // Starting from the next key after x.keyi , move each key left
    x.Ctmp2 = x.Ctmp2+1        // Update the node x such that key k is removed from the node
x.n = x.n - 1                 // Update the length of node x
free(z)                       // Free the pointer to node z as it has already been merged
B-TREE-DELETE(y,k)            // Recursively delete key k from y.
DISK-WRITE(x)                 // Write nodes x,y,z to disk
DISK-WRITE(y)
DISK-WRITE(z)
Return true

```

### Case 3

In this case, we are trying to search for k while traversing down the tree. During this process, we review each internal node to check if it has t-1 keys and if it is, we will need to adjust the number of keys so that it is not minimum. As we are traversing down, we also identify root x.c<sub>i</sub> of the subtree that must contain k.

```

x,i = B-Tree-Search(x,k)      //Returns the node and location of key k
x.ci = Find-Root(x,k)         //Returns Root x.ci
DISK-READ(x.ci)
if x.ci.n > t - 1             // If the node contains t - 1 keys, we can continue
    B-TREE-DELETE(x.ci,k)
DISK-READ(x.ci+1)           //

```

### Case 3a

```

Else if (x.ci+1.n) >= t      //Check if the right sibling has t or more keys
    x.ci.keyn+2 = x.keyi    //Push the key from x down to x.ci
    x.ci.n = x.ci.n + 1      //Increment the length of x.ci to reflect the additional key
    x.keyi = x.ci+1.key1    //Move key from x.ci's immediate right sibling up into x
    x.ci.Cn+1 = x.ci+1.C1  //Move the child pointer from the sibling into x.ci
    for tmp1 = 1 to x.ci+1.n

```

```

        x.Ci+1.keyj = x.Ci+1.keyj+1 //Move each key in sibling left.
x.Ci+1.n = x.Ci+1.n - 1 // Reduce the length of the right sibling to reflect the change.
DISK-WRITE(x) // Write all nodes to disk
DISK-WRITE(x.Ci)
DISK-WRITE(x.Ci+1)
B-TREE-DELETE(x.Ci,k) // Recursively delete key k from the Root x.ci

```

### Case 3a continued

```

Else if (x.Ci-1.n) >= t //Check if the left sibling has t or more keys
    x.Ci.keyn+2 = x.keyi //Push the key from x down to x.ci
    x.Ci.n = x.Ci.n + 1 //Increment the length of x.ci to reflect the additional key
    x.keyi = x.Ci-1.key1 //Move key from x.ci's immediate left sibling up into x
    x.Ci.Cn+1 = x.Ci-1.C1 //Move the child pointer from the sibling into x.ci
    for tmp1 = 1 to x.Ci-1.n
        x.Ci-1.keyj = x.Ci-1.keyj+1 //Move each key in sibling left.
    x.Ci-1.n = x.Ci-1.n - 1 // Reduce the length of the left sibling to reflect the change.
    DISK-WRITE(x) // Write all nodes to disk
    DISK-WRITE(x.Ci)
    DISK-WRITE(x.Ci-1)
    B-TREE-DELETE(x.Ci,k) // Recursively delete key k from the Root x.ci

```

### Case 3b

```

Else //if both nodes have t-1 keys
    i = 1 //Starting with the first key
    while x.keyi < k //Scan each entry till we reach key that is smaller than k
        i = i + 1
    y = x.Ci //Identify index i in x to identify location of k so we can get x.Ci
    // to identify preceding child y.
    z = x.Ci+1 //Identify index i+1 in x to identify location of k so we can get
    // x.Ci+1 to identify successor child z.
    y.keyn+1 = k // Insert key k into y
    for tmp1 = 1 to z.n //for each key in z
        y.keyn+1+tmp1 = z.keytmp1 // Append key from z to the end of y
    y.n = y.n + 1 + z.n // Update the size of y to reflect the merged node (y and z)
    for tmp2 = i + 1 to x.n - 1 // Starting from the next key after x.keyi , move each key left
        x.Ctmp2 = x.Ctmp2+1 // Update the node x such that key k is removed from the node
    x.n = x.n - 1 // Update the length of node x
    free(z) // Free the pointer to node z as it has already been merged
    B-TREE-DELETE(y,k) // Recursively delete key k from y.
    DISK-WRITE(x) // Write nodes x,y,z to disk
    DISK-WRITE(y)
    DISK-WRITE(z)
    Return true

```

### Combined code

B-Tree-Delete-Key(x,k)

//Case 2a – Identify if internal node, find predecessor child and if it has only t-1 keys

if(x.leaf = FALSE)

if (!x.leaf) && (x.n > t - 1)

i = 1

while x.key<sub>i</sub> < k

i = i + 1

y = x.C<sub>i</sub>

DISK-READ(y)

if (y.n > t - 1)

k' = Find-Predecessor-Key (k, y)

B-Tree-Delete-Key(y,k')

x.key<sub>i</sub> = k'

Disk-Write(x)

Return true

//Case 2b - Identify successor child and check if successor has t - 1 keys.

else if (x.C<sub>i+1</sub>.n > t-1)

z = x.C<sub>i+1</sub>

DISK-READ(z)

k' = Find- Successor-Key (k, y)

B-Tree-Delete-Key(z,k')

x.key<sub>i</sub> = k'

Disk-Write(x)

Return true

//Case 2c – if preceding as well as successor child has only t-1 keys

else if (z.n == t-1) && (y.n == t-1)

y.key<sub>n+1</sub> = k

for tmp1 = 1 to z.n

y.key<sub>n+1+ tmp1</sub> = z.key<sub>tmp1</sub>

y.n = y.n + 1 + z.n

for tmp2 = i + 1 to x.n -1

x.C<sub>tmp2</sub> = x.C<sub>tmp2+1</sub>

x.n = x.n-1

free(z)

B-TREE-DELETE(y,k)

DISK-WRITE(x)

DISK-WRITE(y)

DISK-WRITE(z)

Return true

// Case 1 (k is in leaf node)

else if(x.leaf == True)

if (x.leaf) && (x.n > t - 1)

```

for i=1 to x.n
  if x.keyi == k
    DELETE-Key(x,k)
    DISK-WRITE(x)
    Return true

```

//Case 3b

```

else
  x,i = B-Tree-Search(x,k)
  x.ci = Find-Root(x,k)
  DISK-READ(x.ci)
  if x.ci.n > t - 1
    B-TREE-DELETE(x.ci,k)
  DISK-READ(x.ci+1)
  if (x.ci+1.n) >=t
    x.ci.keyn+2 = x.keyi
    x.ci.n = x.ci.n + 1
    x.keyi = x.ci+1.key1
    x.ci.cn+1 = x.ci+1.c1
    for tmp1 = 1 to x.ci+1.n
      x.ci+1.keyj = x.ci+1.keyj+1
    x.ci+1.n = x.ci+1.n - 1
    DISK-WRITE(x)
    DISK-WRITE(x.ci)
    DISK-WRITE(x.ci+1)
    B-TREE-DELETE(x.ci,k)
  Else if (x.ci-1.n) >=t
    x.ci.keyn+2 = x.keyi
    x.ci.n = x.ci.n + 1
    x.keyi = x.ci-1.key1
    x.ci.cn+1 = x.ci-1.c1
    for tmp1 = 1 to x.ci-1.n
      x.ci-1.keyj = x.ci-1.keyj+1
    x.ci-1.n = x.ci-1.n - 1
    DISK-WRITE(x)
    DISK-WRITE(x.ci)
    DISK-WRITE(x.ci-1)
    B-TREE-DELETE(x.ci,k)

```

//Case 3b

```

Else
  i = 1
  while x.keyi < k
    i = i + 1
  y = x.ci
  z = x.ci+1

```

```

y.keyn+1 = k
for tmp1 = 1 to z.n
    y.keyn+1+tmp1 = z.keytmp1
y.n = y.n + 1 + z.n
for tmp2 = i + 1 to x.n - 1
    x.Ctmp2 = x.Ctmp2+1
x.n = x.n - 1
free(z)
B-TREE-DELETE(y,k)
DISK-WRITE(x)
DISK-WRITE(y)
DISK-WRITE(z)
Return true

```

In summary, during the traversal of the tree, we identify, at each node the criteria on where it is an internal node and how many keys exist in this node and possibly its siblings and then ensure that the rules of the structure of B-tree are maintained at all times. At times, we need to adjust keys, either by rearranging the node's children or by backing up and then adjusting the keys.

**Complexity analysis:** To delete a key from the B-Tree, we have to use the search operation and then adjust the node or siblings. Due to this, the worst-case scenario for delete operation in B-Tree is  $O(\log n)$  and the average case is  $\Theta(\log n)$ .

## Problem 2 (3 + 1 points) Disjoint Sets

In the depth-determination problem, we maintain a forest  $F = \{T_i\}$  of rooted trees under three operations:

*Make-Tree*( $v$ ) creates a tree whose only node is  $v$ ,

*Find-Depth*( $v$ ) returns the depth of node  $v$ ,

*Graft* ( $r$ ;  $v$ ) makes node  $r$  (the root of a tree) as a child of node  $v$  of a different tree.

**(1) (1 point) Suppose a tree representation similar to a disjoint-set forest is used:  $v.p$  is the parent of node  $v$ , except  $v.p=v$  if  $v$  is a root. Suppose further that we implement *Graft* ( $r$ ;  $v$ ) by setting  $r.p=v$  and *Find-Depth*( $v$ ) by following the find path up to the root and returning a count of all nodes other than  $v$  encountered. Show that the worst-case running time of a sequence of  $m$  *Make-Tree* (total  $n$ ), *Find-Depth* (total  $n$ ), and *Graft* (total  $n$ ) operations is  $\Theta(n^2)$ .**

Answer: If we use disjoint-set data structure, *Make-Tree* takes  $\Theta(1)$  time.

*Graft* is basically a union operation; thus, it takes  $\Theta(1)$  times.

The cost of *Find-Depth* depends on the depth of the given node i.e. linear.

For a sequence of  $n$  operations, the depth of a node is  $O(n)$ , thus, for the worst-case  $T(n) = n \cdot O(n) = O(n^2)$ .

For example, let  $k = m/3$  be an integer, considering a sequence of operations with  $k + 1$  *Make-Trees* creating  $k + 1$  single-node trees,  $k$  *Grafts* forming a single path, and  $k - 1$  *FIND-DEPTH* for the leaf node, then the running time of the  $n$  operations is

$$T(n) = (k + 1) * \Theta(1) + k \Theta(1) + (k - 1) * k = \Omega(m^2).$$

Hence the worst-case running time is  $\Theta(m^2)$ .

- (2) (1 point) Give an algorithm of *Find-Depth* by modifying *Find-Set*. Your algorithm should perform path compression and its running time should be linear in the length of the find path. Make sure your algorithm updates the pseudo-distances correctly.

Answer: According to the definition of  $v.d$  that the sum of the pseudo-distances along the path from  $v$  to root of its set  $S_i$  equals to the depth of  $v$  in  $T_i$ , *Find-Depth* can be implemented by modifying *Find-Set* in such a way: assume the path is composed of  $v_0, \dots, v_k$  where  $v_k$  is the root, for every node  $v_i$  along the path, update depth of  $v$  in  $T_i$  is  $\sum_{j=0}^k v_j.d$ , i.e., with path compression, whenever the parent pointer of a node changes, the pseudo-distances is updated by the sum of its ancestor's pseudo-distances.

FIND-SET ( $v$ )

1. If ( $v \neq v.p$ ) then
2. ( $v.p, d$ ) = FIND-SET( $v.p$ )
3.  $v.d = v.d + d$
4. return ( $v.p, v.d$ )
5. else
6. return ( $v, 0$ )
7. end if

- (3) (1 point) Give an algorithm of *Graft* ( $r, v$ ), which combines the sets containing  $r$  and  $v$ , by modifying the *Union* and *Link* procedures. Make sure that your algorithm updates pseudo-distances correctly. Note that the root of a set  $S_i$  is not necessarily the root of the corresponding tree  $T_i$ .

Answer: To implement *Graft* we need to find  $v$ 's actual depth and add it to the pseudodistance of the root of the tree  $S_i$  which contains  $r$ .

GRAFT ( $r, v$ )

1. ( $X, d1$ ) = Find-set( $r$ )                      // find the parent of set
2. ( $y, d2$ ) = Find-set( $v$ )
3. if ( $x.rank > y.rank$ ) then                    //check whose rank is higher
4.      $y.p = x$
5.      $x.d = x.d + d2 + y.d$
6. Else
7.      $x.p = y$
8.      $x.d = x.d + d2$
9.     if ( $x.rank == y.rank$ ) then
10.          $y.rank = y.rank + 1$
11.     end if
12. end if

- (4) (1 bonus point) Give a tight bound on the worst-case running time of a sequence of  $m$  *Make-Tree*, *Find-Depth*, and *Graft* operations,  $n$  of which are *Make-Tree* operations.

Answer: The three implemented operations have the same asymptotic running time as MAKE, FIND, and UNION for disjoint sets, so the worst-case runtime of  $m$  such operations,  $n$  of which are MAKE-TREE operations, is  $O(m\alpha(n))$ .